


Un proceso es una instancia de un programa que está en ejecución

Cada instancia tiene su propio espacio de direcciones y estado de ejecución

Cuando se ejecuta un programa, el sistema de operación copia el módulo ejecutable en una imagen del programa en memoria principal



Unix identifica los procesos por un valor entero (process ID o PID)

Además cada proceso tiene un PID del proceso que lo creó (proceso padre)

Si el padre termina antes que su hijo, éste es adoptado por un proceso del sistema

```
#include <unistd.h>

pid_t getpid(void);
pid_t getppid(void);

printf("My process ID is %ld\n", getpid());
```



El comando ps muestra información de los procesos

Se ejecuta a través de la línea de comando de UNIX

Por defecto muestra información de los procesos del usuario

UID – user ID

PID – process ID

PPID – parent process ID

C – (obsolete)

STIME – starting time of the process

TTY – controlling terminal

TIME – cumulative execution time

CMD – command name

```

UID      PID  PPID  C STIME TTY      TIME CMD
root      1    0    0 Jan08 ?        00:00:01 init [2]
root      2    0    0 Jan08 ?        00:00:00 [kthreadd]
root      3    2    0 Jan08 ?        00:00:01 [ksoftirqd/0]
root      6    2    0 Jan08 ?        00:00:00 [migration/0]
root      7    2    0 Jan08 ?        00:00:00 [watchdog/0]
root      8    2    0 Jan08 ?        00:00:00 [migration/1]
root     10    2    0 Jan08 ?        00:00:00 [ksoftirqd/1]
root     12    2    0 Jan08 ?        00:00:00 [watchdog/1]
    
```



# Creación de Procesos



Un proceso crea otro a través de la llamada al sistema `fork()`

El proceso que invoca el `fork` se conoce como padre y el nuevo proceso es el hijo

El hijo recibe una copia de la imagen en memoria del proceso padre. Cada proceso tiene su propio espacio de direcciones

Ambos procesos continúan con la instrucción que sigue después del `fork`



```
#include <unistd.h>

pid_t fork(void);
```

El valor de retorno del fork se usa para determinar quién es el proceso padre y quién es el hijo

El valor de retorno para el hijo es 0 (cero) mientras que el padre recibe el PID del hijo



```
bitvise xterm - garon.tip - auliyahoc.usb.ve:zz
FORK(2)                                Linux Programmer's Manual                                FORK(2)

NAME
    fork - create a child process

SYNOPSIS
    #include <unistd.h>

    pid_t fork(void);

DESCRIPTION
    fork() creates a new process by duplicating the calling process. The new process, referred to as the child, is an exact duplicate of the calling process, referred to as the parent, except for the following points:

    * The child has its own unique process ID, and this PID does not match the ID of any existing process group (setpgid(2)).

    * The child's parent process ID is the same as the parent's process ID.

    * The child does not inherit its parent's memory locks (mlock(2), mlockall(2)).

    * Process resource utilizations (getrusage(2)) and CPU time counters (times(2)) are reset to zero in the child.

    * The child's set of pending signals is initially empty (sigpending(2)).

    * The child does not inherit semaphore adjustments from its parent (semop(2)).

    * The child does not inherit record locks from its parent (fcntl(2)).
```



- \* The child process is created with a single thread—the one that called `fork()`. The entire virtual address space of the parent is replicated in the child, including the states of mutexes, condition variables, and other pthreads objects; the use of `pthread_atfork(3)` may be helpful for dealing with problems that this can cause.
- \* The child inherits copies of the parent's set of open file descriptors. Each file descriptor in the child refers to the same open file description (see `open(2)`) as the corresponding file descriptor in the parent. This means that the two descriptors share open file status flags, current file offset, and signal-driven I/O attributes (see the description of `F_SETOWN` and `F_SETSIG` in `fcntl(2)`).
- \* The child inherits copies of the parent's set of open message queue descriptors (see `mq_overview(7)`). Each descriptor in the child refers to the same open message queue description as the corresponding descriptor in the parent. This means that the two descriptors share the same flags (mq flags).
- \* The child inherits copies of the parent's set of open directory streams (see `opendir(3)`). POSIX.1-2001 says that the corresponding directory streams in the parent and child may share the directory stream positioning; on Linux/glibc they do not.

#### RETURN VALUE

On success, the PID of the child process is returned in the parent, and 0 is returned in the child. On failure, -1 is returned in the parent, no child process is created, and errno is set appropriately.

#### ERRORS

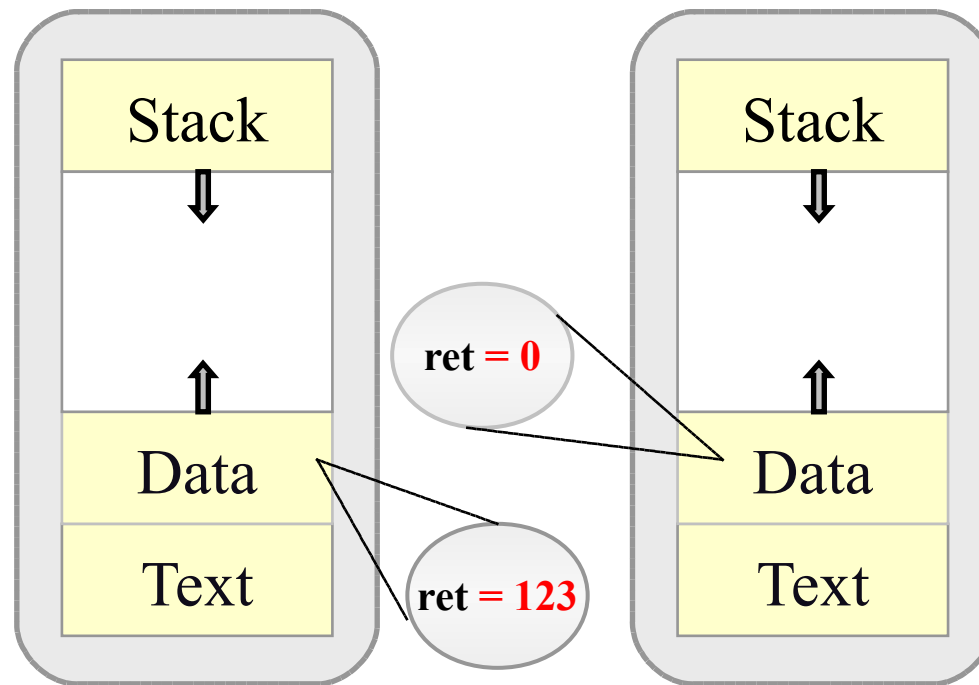
EAGAIN `fork()` cannot allocate sufficient memory to copy the parent's page tables and allocate a task structure for the child.

EAGAIN It was not possible to create a new process because the caller's `RLIMIT_NPROC` resource limit was encountered. To exceed this limit, the process must have either the `CAP_SYS_ADMIN` or the `CAP_SYS_RESOURCE` capability.





`fork()`




El proceso hijo hereda del padre una copia idéntica de su memoria

Registros de CPU

Todos los archivos abiertos

La ejecución prosigue de forma concurrente a partir de la instrucción que sigue el fork

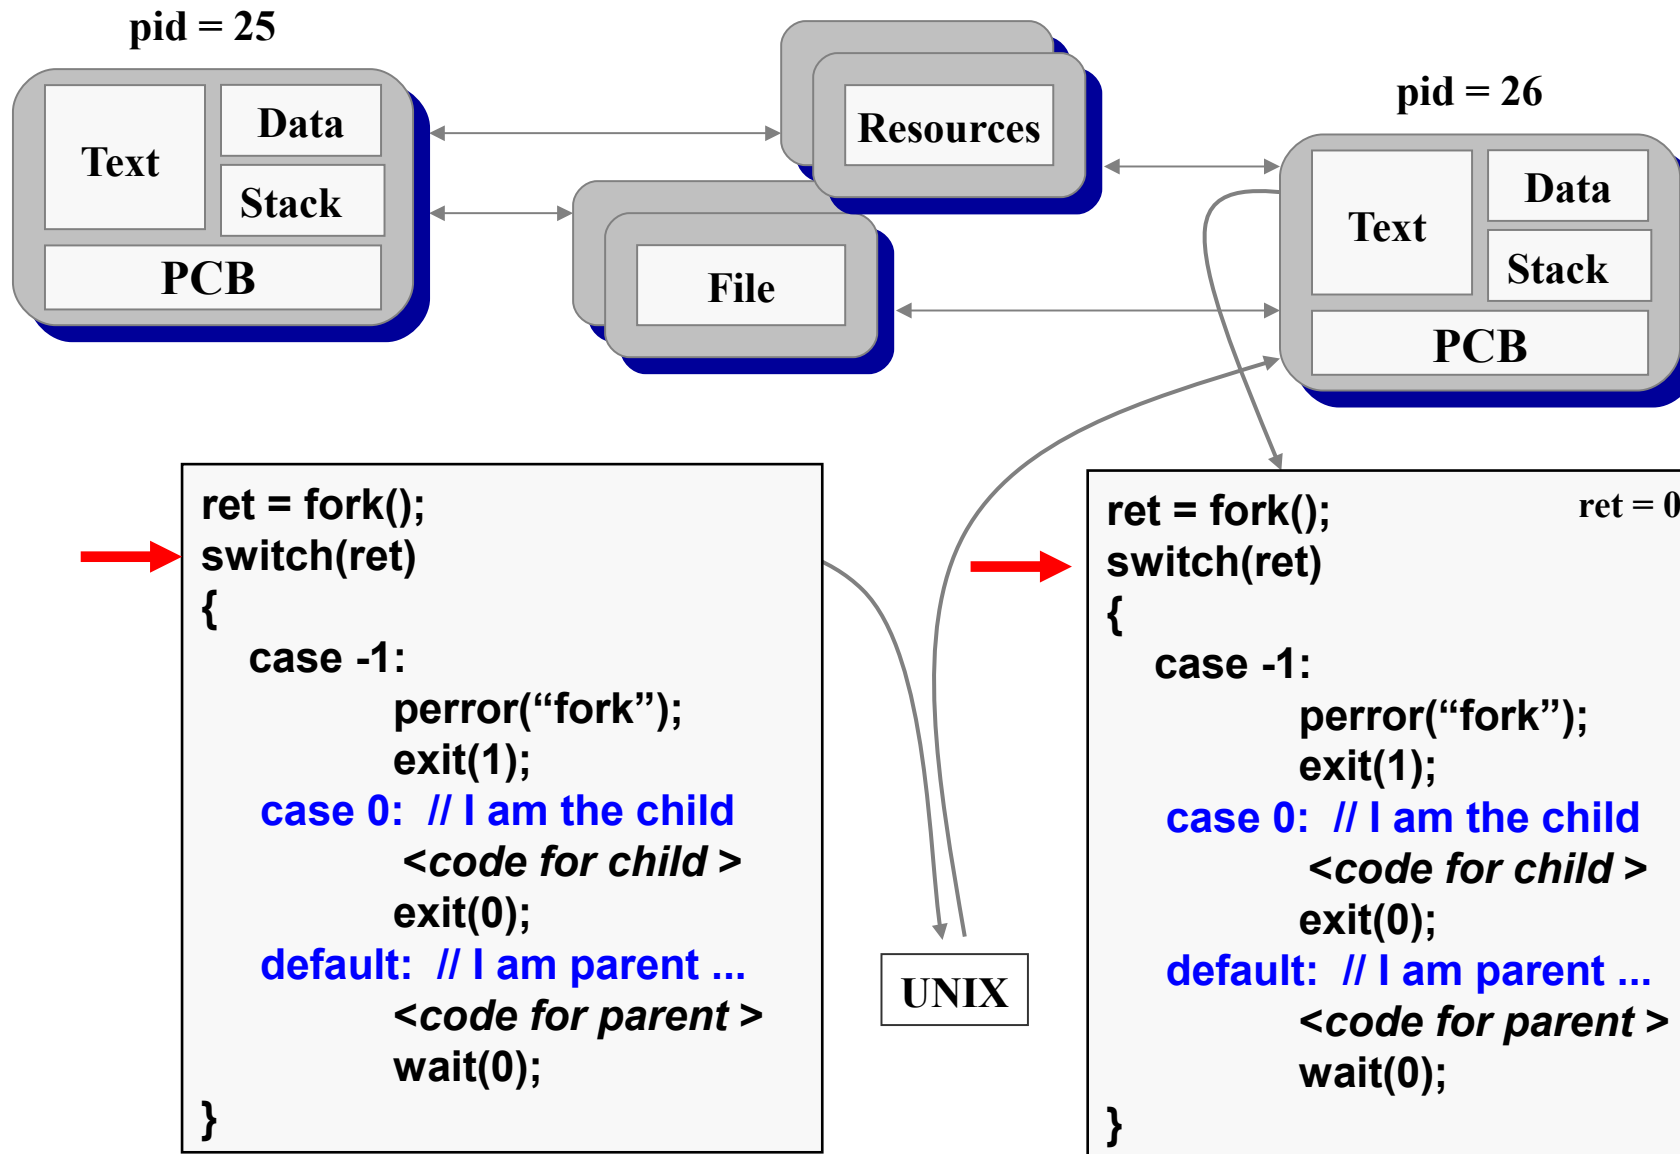


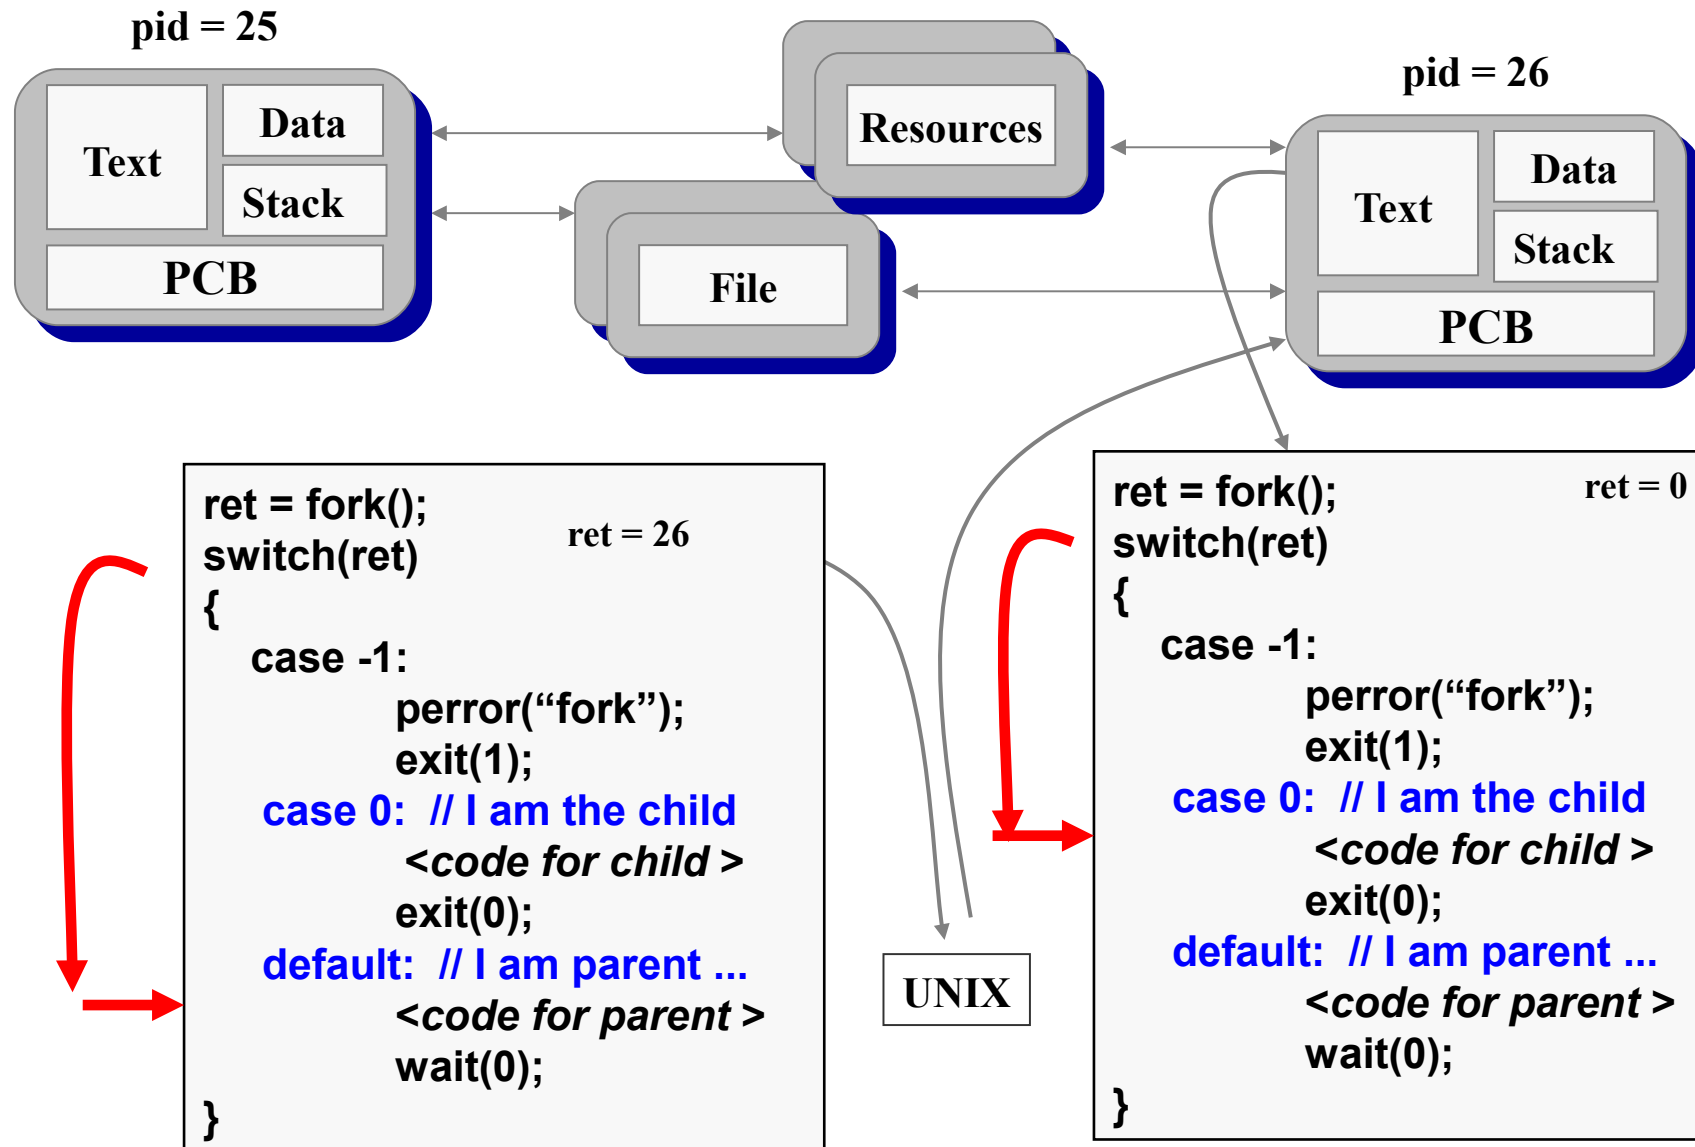


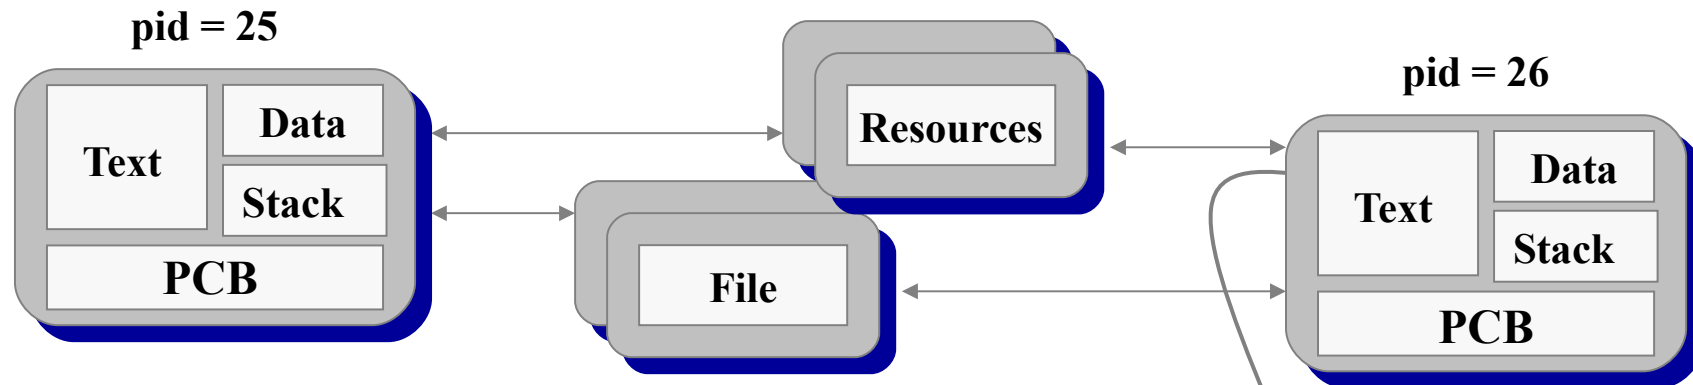
```
→ ret = fork();  
  switch(ret)  
  {  
    case -1:  
      perror("fork");  
      exit(1);  
    case 0: // I am the child  
      <code for child >  
      exit(0);  
    default: // I am parent ...  
      <code for parent >  
      wait(0);  
  }
```

UNIX









```

ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        <code for child >
        exit(0);
    default: // I am parent ...
        <code for parent >
        wait(0);
}

```

ret = 26

UNIX


```

ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        <code for child >
        exit(0);
    default: // I am parent ...
        <code for parent >
        wait(0);
}

```

ret = 0





```

ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        <code for child >
        exit(0);
    default: // I am parent ...
        <code for parent >
        wait(0);
}

```

ret = 26

UNIX


```

ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        <code for child >
        exit(0);
    default: // I am parent ...
        <code for parent >
        wait(0);
}

```

ret = 0





```
ret = fork();
switch(ret)
{
    case -1:
        perror("fork");
        exit(1);
    case 0: // I am the child
        <code for child >
        exit(0);
    default: // I am parent ...
        <code for parent >
        Wait(0);
        <....>
}
```

ret = 26

UNIX





Para finalizar la ejecución, el hijo puede invocar la llamada al sistema `exit(resultado)`

Qué hace un `exit`?



## Qué hace un exit?

Salva el resultado= argumento de la llamada exit

Cierra todos los archivos abiertos, conexiones

Libera la memoria

Chequea si el padre está vivo

Si el padre vive, mantiene el valor de retorno hasta que el padre lo solicita con un **wait**. En este caso el proceso hijo no muere y entra en un estado de **zombie o defunct**

Si el padre no está vivo, el hijo termina (muere)



```
EXIT(3)                                Linux Programmer's Manual                                EXIT(3)

NAME
    exit - cause normal process termination

SYNOPSIS
    #include <stdlib.h>

    void exit(int status);

DESCRIPTION
    The exit() function causes normal process termination and the value of status & 0377 is returned to the parent (see wait(2)).

    All functions registered with atexit(3) and on_exit(3) are called, in the reverse order of their registration. (It is possible for one of these functions to use atexit(3) or on_exit(3) to register an additional function to be executed during exit processing; the new registration is added to the front of the list of functions that remain to be called.) If one of these functions does not return (e.g., it calls _exit(2), or kills itself with a signal), then none of the remaining functions is called, and further exit processing (in particular, flushing of stdio(3) streams) is abandoned. If a function has been registered multiple times using atexit(3) or on_exit(3), then it is called as many times as it was registered.

    All open stdio(3) streams are flushed and closed. Files created by tmpfile(3) are removed.

    The C standard specifies two constants, EXIT_SUCCESS and EXIT_FAILURE, that may be passed to exit() to indicate successful or unsuccessful termination, respectively.

RETURN VALUE
    The exit() function does not return.

CONFORMING TO
    SVr4, 4.3BSD, POSIX.1-2001, C89, C99.
```



El proceso padre puede querer esperar a que sus procesos hijos finalicen

Se utiliza la llamada al sistema `wait()` para esperar por un proceso hijo

El proceso padre se bloquea hasta que alguno de sus hijos termine. La llamada retorna el PID del proceso que finalizó o -1 si no hay procesos hijos

Para esperar por un proceso en particular se usa `waitpid()`



```
#include <sys/types.h>
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```



```
WAIT(2)                                     Linux Programmer's Manual                               WAIT(2)

NAME
    wait, waitpid, waitid - wait for process to change state

SYNOPSIS
    #include <sys/types.h>
    #include <sys/wait.h>

    pid_t wait(int *status);

    pid_t waitpid(pid_t pid, int *status, int options);

    int waitid(idtype_t idtype, id_t id, siginfo_t *info, int options);

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):

    waitid():
        _SVID_SOURCE || _XOPEN_SOURCE >= 500 || _XOPEN_SOURCE && _XOPEN_SOURCE_EXTENDED
        || /* Since glibc 2.12: */ _POSIX_C_SOURCE >= 200809L

DESCRIPTION
    All of these system calls are used to wait for state changes in a child of the calling
    process, and obtain information about the child whose state has changed. A state change is
    considered to be: the child terminated; the child was stopped by a signal; or the child was
    resumed by a signal. In the case of a terminated child, performing a wait allows the system
    to release the resources associated with the child; if a wait is not performed, then the ter-
    minated child remains in a "zombie" state (see NOTES below).
```





If a child has already changed state, then these calls return immediately. Otherwise they block until either a child changes state or a signal handler interrupts the call (assuming that system calls are not automatically restarted using the SA\_RESTART flag of `sigaction(2)`). In the remainder of this page, a child whose state has changed and which has not yet been waited upon by one of these system calls is termed waitable.

#### `wait()` and `waitpid()`

The `wait()` system call suspends execution of the calling process until one of its children terminates. The call `wait(&status)` is equivalent to:

```
waitpid(-1, &status, 0);
```

The `waitpid()` system call suspends execution of the calling process until a child specified by pid argument has changed state. By default, `waitpid()` waits only for terminated children, but this behavior is modifiable via the options argument, as described below.

The value of pid can be:

- < -1 meaning wait for any child process whose process group ID is equal to the absolute value of pid.
- 1 meaning wait for any child process.
- 0 meaning wait for any child process whose process group ID is equal to that of the calling process.
- > 0 meaning wait for the child whose process ID is equal to the value of pid.



If status is not NULL, wait() and waitpid() store status information in the int to which it points. This integer can be inspected with the following macros (which take the integer itself as an argument, not a pointer to it, as is done in wait() and waitpid(!):

**WIFEXITED(status)**

returns true if the child terminated normally, that is, by calling exit(3) or \_exit(2), or by returning from main().

**WEXITSTATUS(status)**

returns the exit status of the child. This consists of the least significant 8 bits of the status argument that the child specified in a call to exit(3) or \_exit(2) or as the argument for a return statement in main(). This macro should only be employed if WIFEXITED returned true.

**WIFSIGNALED(status)**

returns true if the child process was terminated by a signal.

**WTERMSIG(status)**

returns the number of the signal that caused the child process to terminate. This macro should only be employed if WIFSIGNALED returned true.

**WCOREDUMP(status)**

returns true if the child produced a core dump. This macro should only be employed if WIFSIGNALED returned true. This macro is not specified in POSIX.1-2001 and is not available on some UNIX implementations (e.g., AIX, SunOS). Only use this enclosed in #ifdef WCOREDUMP ... #endif.

**WIFSTOPPED(status)**

returns true if the child process was stopped by delivery of a signal; this is only





**getpid** retorna el identificador del proceso

```
pid = getpid();
```

**getppid** retorna el identificador del padre

Procesos Zombies pueden ser vistos usando el comando **ps** .

En el caso de procesos zombies aparecerá **<defunct>** en la columna de comando




```
int main()
{
    pid_t pid;
    int x = 1;


    pid = fork();
    if (pid != 0) {
        printf("parent: x = %d\n", --x);
        exit(0);
    } else {
        printf("child: x = %d\n", ++x);
        exit(0);
    }
}
```



```
void fork2()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```



```
void fork3()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("L2\n");  
    fork();  
    printf("Bye\n");  
}
```



```
void fork4()  
{  
    printf("L0\n");  
    if (fork() != 0) {  
        printf("L1\n");  
        if (fork() != 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



```
void fork5()  
{  
    printf("L0\n");  
    if (fork() == 0) {  
        printf("L1\n");  
        if (fork() == 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



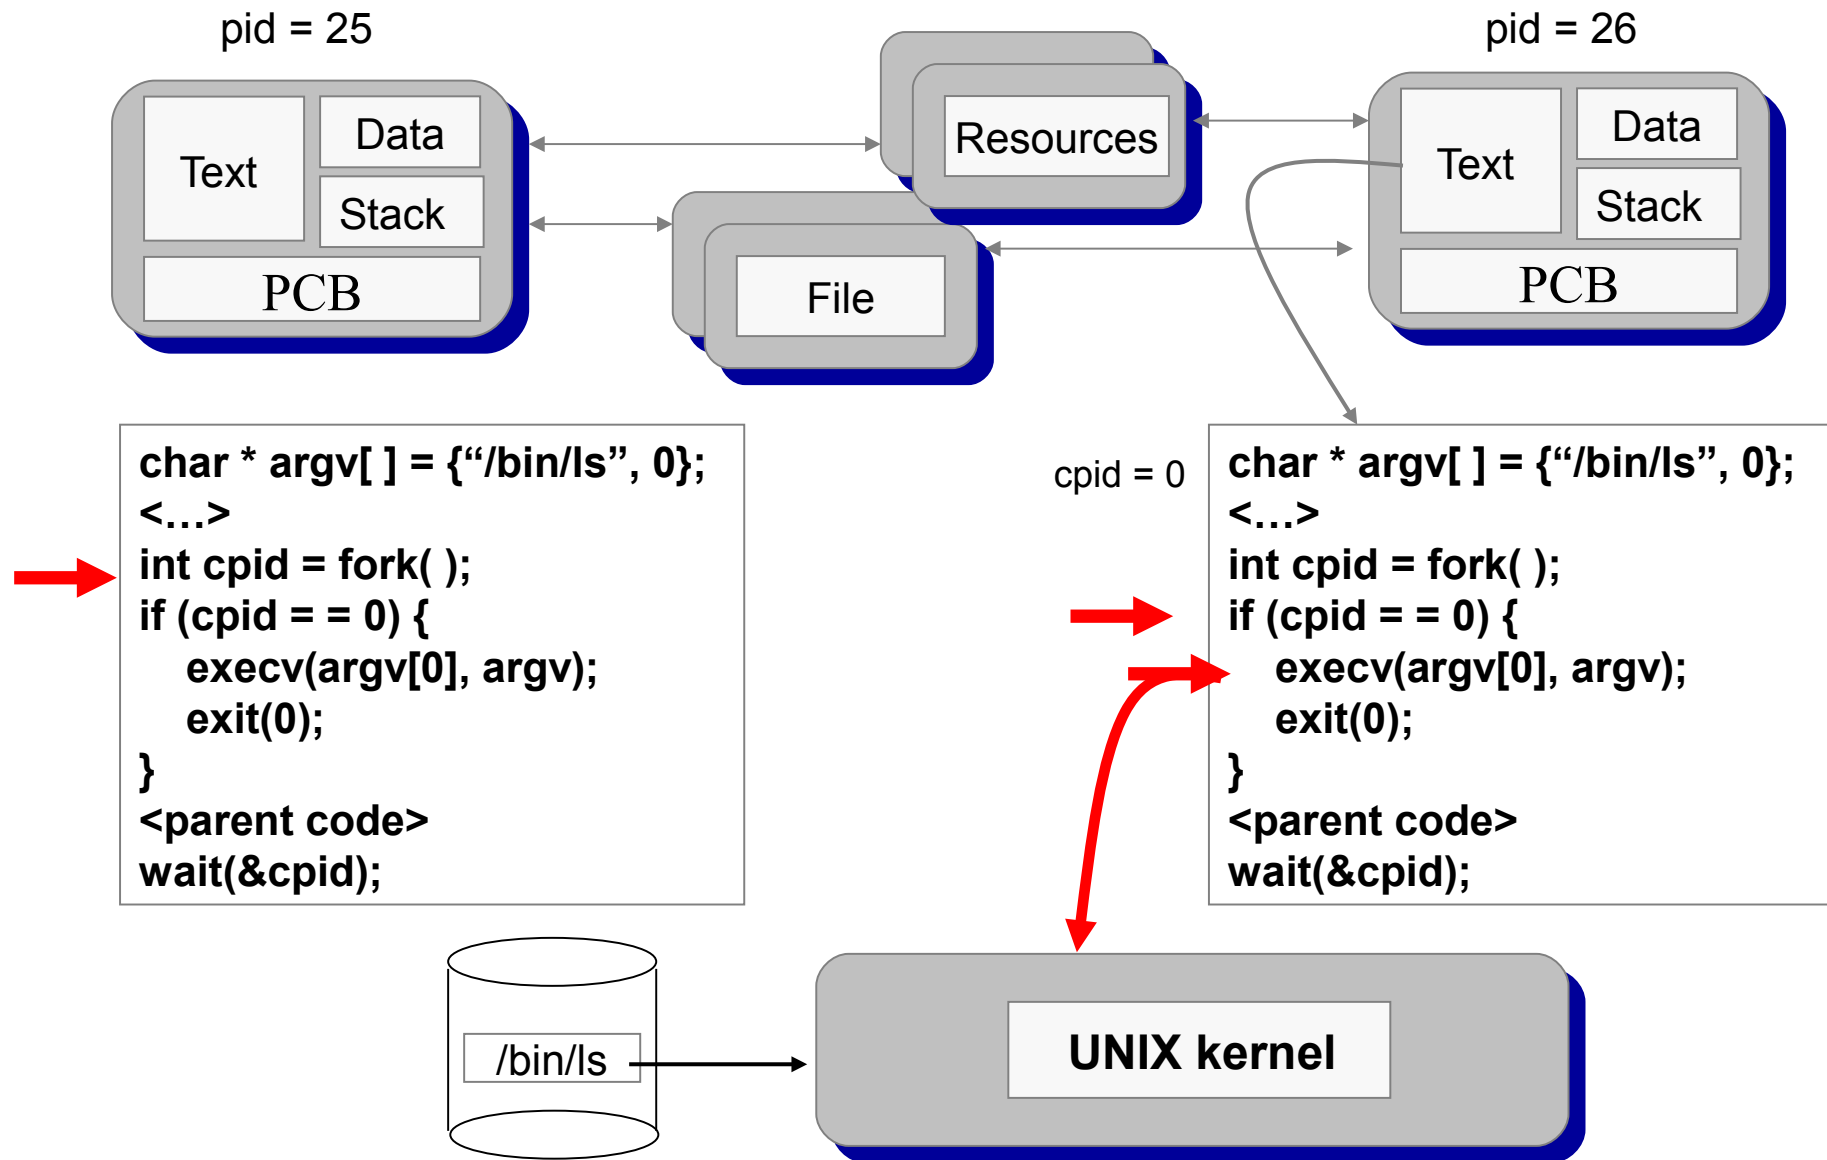
**execv** ejecuta un archivo transformando el proceso llamador en un nuevo proceso . Después de la ejecución correcta de **execv** no hay retorno al proceso llamador

**execv(const char \* path, char \* const argv[])**


**path** camino completo al archivo a ser ejecutado

**argv** arreglo de argumentos para el programa a ejecutar. Cada argumento es una cadena de caracteres terminado con el caracter nulo. El primer argumento es el nombre del programa y el último es NULL











```
char * argv[ ] = {"/bin/ls", 0};
<...>
int cpid = fork( );
if (cpid == 0) {
    execv(argv[0], argv);
    exit(0);
}
<parent code>
wait(&cpid);
```



Exec destruye la imagen del proceso llamador. Una nueva imagen es construida a partir del ejecutable ls





```
char * argv[ ] = {"/bin/lS", 0};
<...>
int cpid = fork( );
if (cpid == 0) {
    execv(argv[0], argv);
    exit(0);
}
<parent code>
wait(&cpid);
```

```
<first line of ls>
<...>
<...>
<...>
exit(0);
```



```
#include <stdio.h>
#include <unistd.h>

char * argv[] = {"/bin/ls", "-l", 0};
int main()
{
    int pid, status;

    if ( (pid = fork() ) < 0 )
    {
        printf("Fork error \n");
        exit(1);
    }
    if(pid == 0) { /* Child executes here */
        execv(argv[0], argv);
        printf("Exec error \n");
        exit(1);
    } else /* Parent executes here */
        wait(&status);
    printf("Hello there! \n");
    return 0;
}
```



Execl similar a execv pero los argumentos para el nuevo programa se pasan como una lista y no un vector

```
execl("/bin/ls", "/bin/ls", "-l", 0);
```

es equivalente a

```
char * argv[] = {" /bin/ls", "-l", 0};  
execv(argv[0], argv);
```

execl es usado principalmente cuando se conoce el número de argumentos a ser pasados



```
int childPid;
char * const argv[ ] = {...};

main {
    childPid = fork();
    if(childPid == 0)
    {
        // I am child ...
        // Do some cleaning, close files
        execv(argv[0], argv);
    }
    else
    {
        // I am parent ...
        <code for parent process>
        wait(0);
    }
}
```



```
EXEC(3)                                Linux Programmer's Manual                                EXEC(3)

NAME
    execl, execlp, execl, execv, execvp, execvpe - execute a file

SYNOPSIS
    #include <unistd.h>

    extern char **environ;

    int execl(const char *path, const char *arg, ...);
    int execlp(const char *file, const char *arg, ...);
    int execl(const char *path, const char *arg,
        ..., char * const envp[]);
    int execv(const char *path, char *const argv[]);
    int execvp(const char *file, char *const argv[]);
    int execvpe(const char *file, char *const argv[],
        char *const envp[]);

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):

    execvpe(): _GNU_SOURCE
```



## DESCRIPTION

The `exec()` family of functions replaces the current process image with a new process image. The functions described in this manual page are front-ends for `execve(2)`. (See the manual page for `execve(2)` for further details about the replacement of the current process image.)

The initial argument for these functions is the name of a file that is to be executed.

The `const char *arg` and subsequent ellipses in the `execl()`, `execlp()`, and `execle()` functions can be thought of as `arg0`, `arg1`, ..., `argn`. Together they describe a list of one or more pointers to null-terminated strings that represent the argument list available to the executed program. The first argument, by convention, should point to the filename associated with the file being executed. The list of arguments must be terminated by a NULL pointer, and, since these are variadic functions, this pointer must be cast `(char *) NULL`.

The `execv()`, `execvp()`, and `execvpe()` functions provide an array of pointers to null-terminated strings that represent the argument list available to the new program. The first argument, by convention, should point to the filename associated with the file being executed. The array of pointers must be terminated by a NULL pointer.

The `execle()` and `execvpe()` functions allow the caller to specify the environment of the executed program via the argument `envp`. The `envp` argument is an array of pointers to null-terminated strings and must be terminated by a NULL pointer. The other functions take the environment for the new process image from the external variable `environ` in the calling process.

#### Special semantics for `execlp()` and `execvp()`

The `execlp()`, `execvp()`, and `execvpe()` functions duplicate the actions of the shell in searching for an executable file if the specified filename does not contain a slash (/) character. The file is sought in the colon-separated list of directory pathnames specified in the `PATH` environment variable. If this variable isn't defined, the path list defaults to the current



```
int main (int argc, char *argv[])
{
pid_t childpid = 0;
int i, nbrOfProcesses;

if (argc != 2)
{
    /* Check for valid number of command-line arguments */
    fprintf(stderr, "Usage: %s <processes>\n", argv[0]);
    return 1;
} // End if

nbrOfProcesses = atoi(argv[1]); // Convert string to an integer
for (i = 1; i < nbrOfProcesses; i++)
{
    childpid = fork();
    if (childpid == -1)
    {
        perror("Fork failed");
        exit(1);
    } // End if
    else if (childpid == 0) // The child
    {
        printf("i:%d process ID: %4ld parent ID: %4ld child ID: %4ld\n",
            i, getpid(), getppid(), childpid);
        sleep(2); // Sleep two seconds
        exit(0);
    } // End if
    else // The parent
        continue;
} // End for


printf("i:%d process ID: %4ld parent ID: %4ld child ID: %4ld\n",
    i, getpid(), getppid(), childpid);

return 0;
} // End main
```





```
% a.out 4
i:1 process ID: 2736 parent ID: 120 child ID: 0
i:2 process ID: 3488 parent ID: 120 child ID: 0
i:4 process ID: 120 parent ID: 40 child ID: 512
i:3 process ID: 512 parent ID: 120 child ID: 0
```



```
int main (int argc, char *argv[])
{
    pid_t childpid = 0;
    int i, nbrOfProcesses;

    if (argc != 2)
    {
        /* Check for valid number of command-line arguments */
        fprintf(stderr, "Usage: %s <processes>\n", argv[0]);
        return 1;
    } // End if

    nbrOfProcesses = atoi(argv[1]); // Convert character string to integer
    for (i = 1; i < nbrOfProcesses; i++)
    {
        childpid = fork();
        if (childpid == -1)
        {
            perror("Fork failed");
            exit(1);
        } // End if
        else if (childpid != 0) // True for a parent
            break;
    } // End for

    // Each parent prints this line
    fprintf(stderr, "i: %d process ID: %4ld parent ID: %4ld child ID: %4ld\n",
            i, (long)getpid(), (long)getppid(), (long)childpid);
    sleep(5); // Sleep five seconds
    return 0;
} // End main
```



```
% a.out 4
```

```
i: 1  process ID:  496  parent ID:   40  child ID: 3232
i: 2  process ID: 3232  parent ID:  496  child ID:  320
i: 3  process ID:  320  parent ID: 3232  child ID: 2744
i: 4  process ID: 2744  parent ID:  320  child ID:    0
```

